LAYCAN

PRE-SEED · DRY BULK, INDIA-BOUND · SEPTEMBER 2026

A freight decision engine for bulk cargo importers.

Freight stops being a daily purchase and becomes a managed position.

THE PROBLEM

Indian steel, power and cement producers import coal and ore by the shipload and buy the ocean freight one voyage at a time, on the day the plant asks for it. Four things are missing, and each of them costs money. There is no threshold that tells a buyer whether today's offer is good relative to waiting, so the decision defaults to *the plant needs coal, fix it*. The vessel class is chosen by habit rather than by computing what the discharge port's draft limit actually allows a ship to load. Nothing is compared afterwards against what could have been paid, so there is no institutional memory. And the freight price risk is carried entirely unhedged — by buyers whose counterparty across every negotiation hedges routinely.

WHAT WE SELL

One page a day, per cargo: a decision memo that commits to an action.

- **Fix or wait**, expressed as a reservation rate — a number today's offer is measured against, not an arrow on a chart
- **Which vessel class, to which port**, with draft-limited cargo intake computed from trim and density, not read off a table
- **Spot, trip charter or contract of affreightment**, chosen on an expected-cost versus tail-risk frontier and capped by what the buyer can actually lift
- **How many FFA lots to sell**, with hedge effectiveness and residual basis risk both stated

We do not claim to forecast the market, and we say so on stage. Dry bulk rates are close to a random walk and a liquid forward market already prices them better than we could. The value sits in the stopping rule, the physical feasibility constraints and the hedge — none of which require forecast superiority. We prove it the only way that counts: by backtesting *decisions* in dollars per tonne against a naive buyer and against perfect hindsight, and publishing the worst quarter next to the average.

WHY IT IS DEFENSIBLE

Two assets compound, and neither of them is software. The first is the **decision-outcome record**: every recommendation stored with its model version and joined to what actually happened — the rate fixed, the days waited, the demurrage paid. Nobody holds this for India-bound bulk, because nobody else is in the decision loop. The second is the **benchmark**. Authoritative daily assessments exist for the world's liquid routes; they do not exist for Newcastle to Paradip. A platform sitting inside enough Indian importers' fixture flow can construct the reference nobody publishes, and index businesses become infrastructure.

There is also a structural reason the incumbents are unlikely to follow us here: their customers are shipowners, operators and brokers. A product whose job is to tell a charterer to wait is adversarial to the other side of their book.

STATUS

Pre-product. A six-person team on a five-week build for SIH 2026 — two ML, two full-stack, one data, one design. Design-partner access to a national steel producer's live procurement problem through the sponsoring ministry.

MARKET

- India is among the world's largest importers of both coking and thermal coal, at a scale of tens of millions of tonnes of coking coal a year
- Global seaborne dry bulk trade is of the order of five billion tonnes a year
- Ocean freight is commonly a material double-digit percentage of the delivered cost of imported coal, and it swings with the rate cycle
- Our wedge: Indian importers moving roughly 0.5 to 10 Mtpa with no chartering desk — big enough for the savings to matter, small enough never to have built the capability

All four figures are order-of-magnitude estimates and are flagged for verification before any of them is quoted to a counterparty.

BUSINESS MODEL

An annual enterprise subscription, tiered by covered tonnage and anchored to measured savings on that tonnage rather than to seats. We are deliberately not asserting a price yet: it should be set against the first shadow-mode result, not against a spreadsheet.

THE ASK

Indicatively **$600K to $1M for eighteen months**. The largest single line item is a licensed freight index feed and satellite AIS coverage — the one purchase that converts our biggest acknowledged weakness into an operating advantage. Then two quantitative hires and two paid shadow-mode deployments. The final number should be set from real licence quotes, not from a market multiple.

READ PAGE 2 BEFORE YOU QUOTE ANY NUMBER ON PAGE 1.
Written without live data access. Page 2 states which claims are verified, which are estimates, and how to falsify us.

Page 2 — the asterisks

Every startup one-pager has these. Most leave them off.

WHAT IS VERIFIED, AND WHAT IS NOT

This material was produced without live internet access, so **every quantitative and competitive claim in it is an estimate from domain knowledge, not a sourced figure** — port drafts, handling rates, import volumes, freight share of landed cost, competitor pricing, market size. The engineering design does not depend on any of them; the credibility of the pitch does. The repository therefore carries a blocking checklist of sixteen verification items to be closed against primary sources before the deck is shown or a customer approached — about two days of work, and the highest-value two days on the project. An investor who catches one wrong port draft is entitled to discount everything else we say. In that checklist, *unknown* is an acceptable answer and a guess is not.

COMPETITIVE LANDSCAPE – AND THE CLAIM WE MOST NEED TO DISPROVE

The maritime software market is not empty and it would be dishonest to imply otherwise. It divides roughly four ways. **Vessel-tracking and commodity-flow analytics** sell AIS-derived positions, port calls and trade-flow estimates to trading desks, owners and analysts. **Chartering intelligence platforms** sell market data, fixture history and cargo matching to people who charter professionally. **Voyage and commercial management systems** are the system of record for owners running fleets. **Risk and compliance** products sell sanctions and behavioural screening to governments, insurers and banks. Alongside them sit performance-optimisation vendors, broking-desk email tools, port cost specialists, and the index provider the derivatives settle against.

Our central commercial hypothesis is that **none of them sells a prescriptive, dated fix-or-wait threshold combined with a port-feasible vessel choice, an instrument mix and a hedge size, to an industrial importer with no chartering desk**. They sell data and workflow, by the seat, to people who already know what the market is worth. Different product, different buyer. It is also exactly the kind of claim founders talk themselves into, so it sits in the checklist as an item to be attacked rather than confirmed — including a search for Indian freight-tech entrants and for any FFA broker shipping software to physical buyers rather than broking to them.

RISKS, AND WHAT WE ACTUALLY DO ABOUT THEM

The timing edge may not survive a real backtest.

Then we say so, and the pitch changes rather than the honesty. The feasibility and hedging layers create value with no forecasting content at all: a vessel that cannot berth is a hard error regardless of the rate, and an unhedged position is an exposure regardless of direction.

The best rate data is licensed, and we will not scrape it.

Until a licence exists the daily series is a calibrated stochastic process, every simulated figure carries a visible badge, and the simulator is validated against free observables. A licence is a config change and a funding line, not a rewrite.

Public sector procurement is slow and price-led.

So the sponsoring relationship is design partnership and credibility, not first revenue. Initial paid deployments target private importers, where the cycle is shorter.

An incumbent could add a recommendation module.

They would need the decision-outcome dataset and the India-specific port physics — and they would be selling advice against the interests of the owners and brokers who pay them today.

A confident wrong number destroys trust faster than a missing one.

The language models are forbidden from emitting a numeral at all, enforced by a schema validator and a build check rather than by convention. Every figure originates in a versioned, unit-tested solver and carries provenance to its source, licence and timestamp. A critic agent attacks each recommendation before it ships and escalates to a human when it cannot get comfortable.

THE NEXT NINETY DAYS

- **Days 1–14.** Close the verification checklist against primary sources, publish the citation index, and correct or delete every figure that does not survive it.
- **Days 15–45.** Ship the deterministic core — port physics, voyage economics, stopping policy, assignment program — with the decision backtest and a leakage test running in CI.
- **Days 46–70.** Complete the agent layer, the critic, and the provenanced decision memo. Freeze a demo snapshot so nothing depends on venue wifi.
- **Days 71–90.** Take the measured result, whatever it says, to three importers and ask one to run us in shadow mode for a quarter.

HOW TO FALSIFY US

Ask for the capture ratio — the share of theoretically available timing value the policy actually captured, measured walk-forward and net of commissions and frictions. Then ask for the worst quarter in the sample, and for the result against the forward curve rather than against a naive buyer. If we cannot produce all three, we have not earned the meeting.